

# Software Development Practical

---

Computer Vision & Deep Learning





# Deep Learning

---

Neurons



# Neuron


Assume we have a 2D data point with  $x_1$  and  $x_2$  coordinate.

Three things happen in a neuron.

1. Multiply each input by a weight

$$x_1 \rightarrow x_1 * w_1$$

$$x_2 \rightarrow x_2 * w_2$$






# Neuron

Assume we have a 2D data point with  $x_1$  and  $x_2$  coordinate.

Three things happen in a neuron.

1. Multiply each input by a weight
2. Add weighted inputs and bias

$$(x_1 * w_1) + (x_2 * w_2) + b$$






# Neuron

Assume we have a 2D data point with  $x_1$  and  $x_2$  coordinate.

Three things happen in a neuron.

1. Multiply each input by a weight
2. Add weighted inputs and bias
3. Pass sum through activation function

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$





# Activation Function

⇒ non-linearity

important ones:


- ReLU →  $[0, \max(x)]$
- Sigmoid →  $[0, 1]$
- Softmax →  $[0, 1]$  & sum = 1

$$s(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

## Activation Functions


**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**tanh**

$$\tanh(x)$$




**ReLU**

$$\max(0, x)$$



**Leaky ReLU**

$$\max(0.1x, x)$$




**Maxout**

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

**ELU**

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



<https://medium.com/@shrutijadon/survey-on-activation-functions-for-deep-learning-9689331ba092>

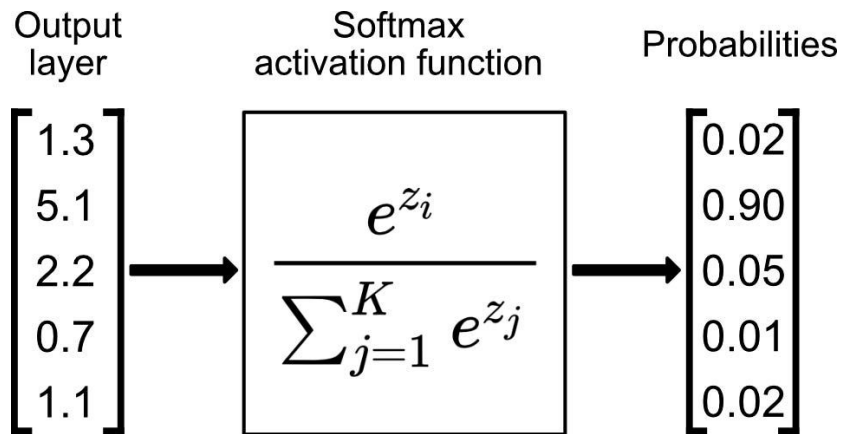


# Activation Function

## Softmax

⇒ convert a vector of real-valued functions into probabilities

⇒ important for e.g. classification






# Example

Three things happen in a neuron.

1. Multiply each input by a weight
2. Add weighted inputs and bias
3. Pass sum through activation function

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$



We have our weight vector and bias:

$$\mathbf{w} = [0, 1] \Rightarrow w_1 = 0; w_2 = 1$$

$$\mathbf{b} = 4$$

We have our input:

$$\mathbf{x} = [2, 3]$$

With the dot product we can write

$$(\mathbf{w} \cdot \mathbf{x}) + b = ((w_1 * x_1) + (w_2 * x_2)) + b$$






# Example

Three things happen in a neuron.

1. Multiply each input by a weight
2. Add weighted inputs and bias
3. Pass sum through activation function

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$



We have our weight vector and bias:

$$\mathbf{w} = [0, 1] \Rightarrow w_1 = 0; w_2 = 1$$

$$\mathbf{b} = 4$$

We have our input:

$$\mathbf{x} = [2, 3]$$

With the dot product we can write

$$\begin{aligned}(w \cdot x) + b &= ((w_1 * x_1) + (w_2 * x_2)) + b \\ &= 0 * 2 + 1 * 3 + 4 \\ &= 7\end{aligned}$$



# Programming 1

Please code the neuron we just described using python.

You have 20 minutes.

1. Implement the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x} = 1 - \sigma(-x).$$

2. Plot the sigmoid function for the range -6 to 6
3. Implement a Neuron class

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

4. Test it with  $\mathbf{w} = [0, 1]$ ,  $\mathbf{b} = 4$  and  $\mathbf{x} = [2, 3]$

```
import numpy as np

def sigmoid(x):
    # TODO

class Neuron:
    def __init__(self, weights, bias):
        # TODO

    def feedforward(self, inputs):
        # TODO
        # Weight inputs, add bias, & use activation function
```



# Programming 1 - Solution

Please code the neuron we just described using python.

You have 20 minutes.

## 1. Implement the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x} = 1 - \sigma(-x).$$

## 2. Plot the sigmoid function for the range -6 to 6


## 3. Implement a Neuron class

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

## 4. Test it with $\mathbf{w} = [0, 1]$ , $\mathbf{b} = 4$ and $\mathbf{x} = [2, 3]$

```
import numpy as np

def sigmoid(x):
    # Our activation function: f(x) = 1 / (1 + e^(-x))
    return 1 / (1 + np.exp(-x))
```





# Programming 1 - Solution

Please code the neuron we just described using python.

You have 20 minutes.

1. Implement the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x} = 1 - \sigma(-x).$$

2. Plot the sigmoid function for the range -6 to 6

3. Implement a Neuron class

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

4. Test it with  $\mathbf{w} = [0, 1]$ ,  $\mathbf{b} = 4$  and  $\mathbf{x} = [2, 3]$

```
class Neuron:
    def __init__(self, weights, bias):
        self.weights = weights
        self.bias = bias

    def feedforward(self, inputs):
        # Weight inputs, add bias, & use activation function
        total = np.dot(self.weights, inputs) + self.bias
        return sigmoid(total)
```



# Programming 1 - Solution

Please code the neuron we just described using python.

You have 20 minutes.

1. Implement the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x} = 1 - \sigma(-x).$$

2. Plot the sigmoid function for the range -6 to 6
3. Implement a Neuron class

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

4. Test it with **w** = [0, 1], **b** = 4 and **x** = [2, 3]

```
weights = np.array([0, 1]) # w1 = 0, w2 = 1
bias = 4 # b = 4
n = Neuron(weights, bias)

x = np.array([2, 3]) # x1 = 2, x2 = 3
print(n.feedforward(x)) # 0.9990889488055994
```



# Deep Learning

---


Multilayer Perceptrons



# Neural Networks

Neural networks are nothing more than multiple neurons connected together.

Back-propagation of loss is used for computing gradients for optimization.





# Neural Networks

## Example


Assume we have a network with all weights having the same weights  $\mathbf{w} = [0, 1]$  and same bias  $\mathbf{b} = \mathbf{0}$ , let  $h_1, h_2, o_1$  denote the outputs of the neurons they represent.

We call it  $h$  because its a *hidden layer*.

E.g. input  $\mathbf{x} = [2, 3]$ :

$$\begin{aligned} h_1 &= h_2 = f(w \cdot x + b) \\ &= f((0 * 2) + (1 * 3) + 0) \\ &= f(3) \\ &= 0.9526 \end{aligned}$$

$$\begin{aligned} o_1 &= f(w \cdot [h_1, h_2] + b) \\ &= f((0 * h_1) + (1 * h_2) + 0) \\ &= f(0.9526) \\ &= \boxed{0.7216} \end{aligned}$$







# Programming 2

Please code the neural network we just described using python.

The neural network has


- 2 inputs
- a hidden layer with 2 neurons ( $h_1$ ,  $h_2$ )
- an output layer with 1 neuron ( $o_2$ )

Each neuron has the same weights and bias  $\mathbf{w} = [0, 1]$  and  $\mathbf{b} = 0$ .

Test the neural network with  $\mathbf{x} = [2, 3]$

You have 20 minutes.

```
class NeuralNetwork:
    def __init__(self):
        # TODO
    def feedforward(self, x):
        # TODO
```





# Programming 2 - Solution

Please code the neural network we just described using python.

The **neural network** has

- 2 inputs
- a hidden layer with 2 neurons (h1, h2)
- an output layer with 1 neuron (o2)

Each neuron has the same weights and bias

$\mathbf{w} = [0, 1]$  and  $\mathbf{b} = 0$ .

Test the neural network with  $\mathbf{x} = [2, 3]$

```
class NeuralNetwork:
    def __init__(self):
        weights = np.array([0, 1])
        bias = 0

        self.h1 = Neuron(weights, bias)
        self.h2 = Neuron(weights, bias)
        self.o1 = Neuron(weights, bias)

    def feedforward(self, x):
        out_h1 = self.h1.feedforward(x)
        out_h2 = self.h2.feedforward(x)
        out_o1 = self.o1.feedforward(np.array([out_h1, out_h2]))
        return out_o1
```



# Programming 2 - Solution

Please code the neural network we just described using python.

The neural network has

- 2 inputs
- a hidden layer with 2 neurons (h1, h2)
- an output layer with 1 neuron (o2)

Each neuron has the same weights and bias  
 $\mathbf{w} = [0, 1]$  and  $\mathbf{b} = 0$ .

Test the neural network with  $\mathbf{x} = [2, 3]$

```
network = NeuralNetwork()
x = np.array([2, 3])
print(network.feedforward(x)) # 0.7216325609518421
```



# Deep Learning

---

Training a Neural Network



# Machine Learning Terms

## Input

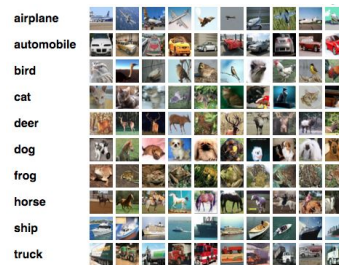
- Data points that we want to process
- e.g. Images, Audio, etc

## Output / Labels


- Corresponding value / label for each datapoint
- e.g. temperature or “car”

## Objective

- Feedback signal
- Measures whether algorithm does a good job



<https://www.analyticsvidhya.com/blog/2020/02/learn-image-classification-cnn-convolutional-neural-networks-3-datasets/>





# Machine Learning Terms

## Input

- Data points that we want to process
- e.g. Images, Audio, etc

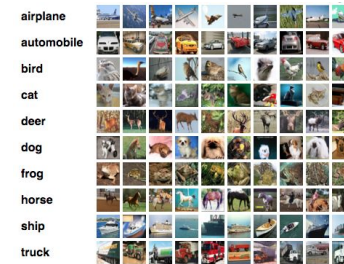
## Output / Labels

- Corresponding value / label for each datapoint
- e.g. temperature or “car”

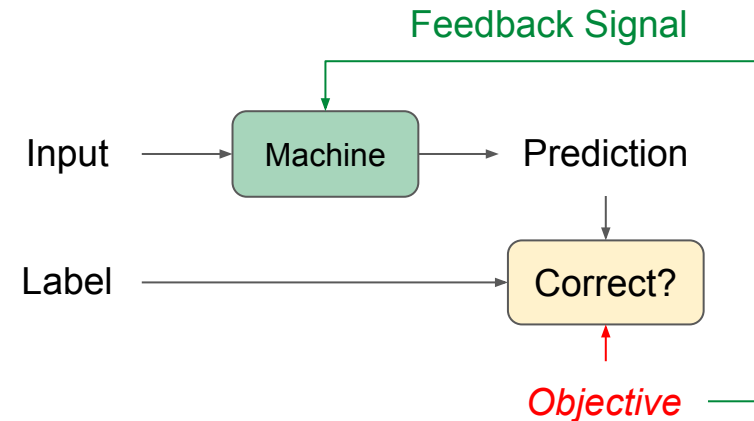
## Objective

- Feedback signal
- Measures whether algorithm does a good job

⇒ Used to change the parameters of our model




<https://www.analyticsvidhya.com/blog/2020/02/learn-image-classification-cnn-convolutional-neural-networks-3-datasets/>




# Objective Function



**data**  
pairs of (x, y) values




**model**  
 $f_{\beta}(x) = y = \beta x$



**cost function / objective**

$$\frac{1}{n} \sum_i^n (y - f_{\beta}(x))^2$$



optimum



# Objective Function / Loss

⇒ Quantifies how “good” our model is


- must be differentiable w.r.t. to parameters
- outputs a scalar
- better predictions = lower loss

## Mean Squared Error

$$MSE(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2$$

Ground truth

Prediction







# Objective Function / Loss


⇒ Quantifies how “good” our model is

## Example for a single sample

Our neural network:  $h(\mathbf{x}; \mathbf{w}, b) = \sigma(\mathbf{w}^T \mathbf{x} + b)$

The loss function: 
$$\begin{aligned} L(\mathbf{x}, y; \mathbf{w}, b) &= (y - \hat{y})^2 \\ &= (y - h(\mathbf{x}; \mathbf{w}, b))^2 \\ &= (y - \sigma(\mathbf{w}^T \mathbf{x} + b))^2 \\ &= (y - \sigma(w_1 \cdot x_1 + w_2 \cdot x_2 + b))^2 \end{aligned}$$

⇒ How much do we need to tweak **w** and **b** to decrease L?





# Partial Derivative


How would  $L$  change if we change  $w_1$ ?

⇒ **partial derivative**  $\frac{\partial L}{\partial w_1}$

How can we get the partial derivative?

$$L(\mathbf{x}, y; \mathbf{w}, b) = (y - \sigma(w_1 \cdot x_1 + w_2 \cdot x_2 + b))^2$$

⇒ **Chain rule**





# Chain Rule

The derivative of the composition of two functions  $f$  and  $g$  can be expressed as

$$h'(x) = f'(g(x)) \cdot g'(x)$$

or equivalently

$$h' = (f \circ g)' = (f' \circ g) \cdot g'.$$

## Example

$$h(x) = (2x^3 + 3x)^2$$

with  $u(x) := 2x^3 + 3x$

$$\begin{aligned} h(u) &= u^2 \\ \frac{\partial h(u)}{\partial u} &= 2u \end{aligned}$$

$$\frac{\partial u(x)}{\partial x} = 6x^2 + 3$$

$$\begin{aligned} \frac{\partial h}{\partial x} &= \frac{\partial h}{\partial u} \cdot \frac{\partial u}{\partial x} \\ &= 2 \cdot (2x^3 + 3x) \cdot (6x^2 + 3) \end{aligned}$$



# Chain Rule

For our example


$$L(\mathbf{x}, y; \mathbf{w}, b) = (y - \sigma(w_1 \cdot x_1 + w_2 \cdot x_2 + b))^2$$

$$g := w_1 \cdot x_1 + w_2 \cdot x_2 + b$$

$$\hat{y} := \sigma(g)$$

$$L := (y - \hat{y})^2$$

$$\begin{aligned} L(\mathbf{x}, y; \mathbf{w}, b) &= (y - \underbrace{\sigma(w_1 \cdot x_1 + w_2 \cdot x_2 + b)}_g)^2 \\ &= (y - \underbrace{\sigma(g)}_{\hat{y}})^2 \\ &= (y - \hat{y})^2 \end{aligned}$$



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial g} \cdot \frac{\partial g}{\partial w_1}$$

$$\frac{\partial L}{\partial \hat{y}} = -2 \cdot (y - \hat{y})$$

$$\frac{\partial \hat{y}}{\partial g} = \sigma(g) \cdot (1 - \sigma(g))$$

$$\frac{\partial g}{\partial w_1} = x_1$$




# Chain Rule

## More complicated example

$$L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$$

$$\frac{\partial L}{\partial w_1} = \boxed{\frac{\partial L}{\partial \hat{y}}} \cdot \frac{\partial \hat{y}}{\partial w_1}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial (1 - \hat{y})^2}{\partial \hat{y}} = -2(y - \hat{y})$$





# Chain Rule

## More complicated example


$$L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \boxed{\frac{\partial \hat{y}}{\partial w_1}}$$

$$\hat{y} = o_1 = f(w_5 \cdot h_1 + w_6 \cdot h_2 + b_3)$$

since  $w_1$  only affects  $h_1$  we can write

$$\boxed{\frac{\partial \hat{y}}{\partial w_1} = \frac{\partial \hat{y}}{\partial h_1} \cdot \frac{\partial h_1}{\partial w_1}} \longrightarrow \boxed{\frac{\partial \hat{y}}{\partial h_1} = w_5 \cdot f'(w_5 \cdot h_1 + w_6 \cdot h_2 + b_3)}$$





# Chain Rule

## More complicated example


$$L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial w_1}$$

$$\hat{y} = o_1 = f(w_5 \cdot h_1 + w_6 \cdot h_2 + b_3)$$

since  $w_1$  only affects  $h_1$  we can write


$$\frac{\partial \hat{y}}{\partial w_1} = \frac{\partial \hat{y}}{\partial h_1} \cdot \frac{\partial h_1}{\partial w_1}$$
$$\frac{\partial \hat{y}}{\partial h_1} = w_5 \cdot f'(w_5 \cdot h_1 + w_6 \cdot h_2 + b_3)$$
$$\frac{\partial h_1}{\partial w_1} = x_1 \cdot f'(w_1 \cdot x_1 + w_2 \cdot x_2 + b_1)$$



# Chain Rule

## More complicated example

$$L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$$



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial h_1} \cdot \frac{\partial h_1}{\partial w_1}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial (1 - \hat{y})^2}{\partial \hat{y}} = -2(y - \hat{y})$$

$$\frac{\partial \hat{y}}{\partial h_1} = w_5 \cdot f'(w_5 \cdot h_1 + w_6 \cdot h_2 + b_3)$$



$$\frac{\partial h_1}{\partial w_1} = x_1 \cdot f'(w_1 \cdot x_1 + w_2 \cdot x_2 + b_1)$$

⇒ So for what is this good now?



# Gradient Descent

1. Initialize parameters at random
2. Compute error
3. Compute gradients w.r.t. parameters
4. Apply update rule
5. Repeat from step 2. until the error does not decrease anymore



[https://www.researchgate.net/figure/Non-convex-optimization-We-utilize-stochastic-gradient-descent-to-find-a-local-optimum\\_fig1\\_325142728](https://www.researchgate.net/figure/Non-convex-optimization-We-utilize-stochastic-gradient-descent-to-find-a-local-optimum_fig1_325142728)

## Update Rule


$$w_1 \leftarrow w_1 - \alpha \frac{\partial L}{\partial w_1}$$

# 3blue1brown

---



Really nice introduction video into neural networks...





# Thanks for your Attention

---

Next Week: Dive deeper into Gradient Descent